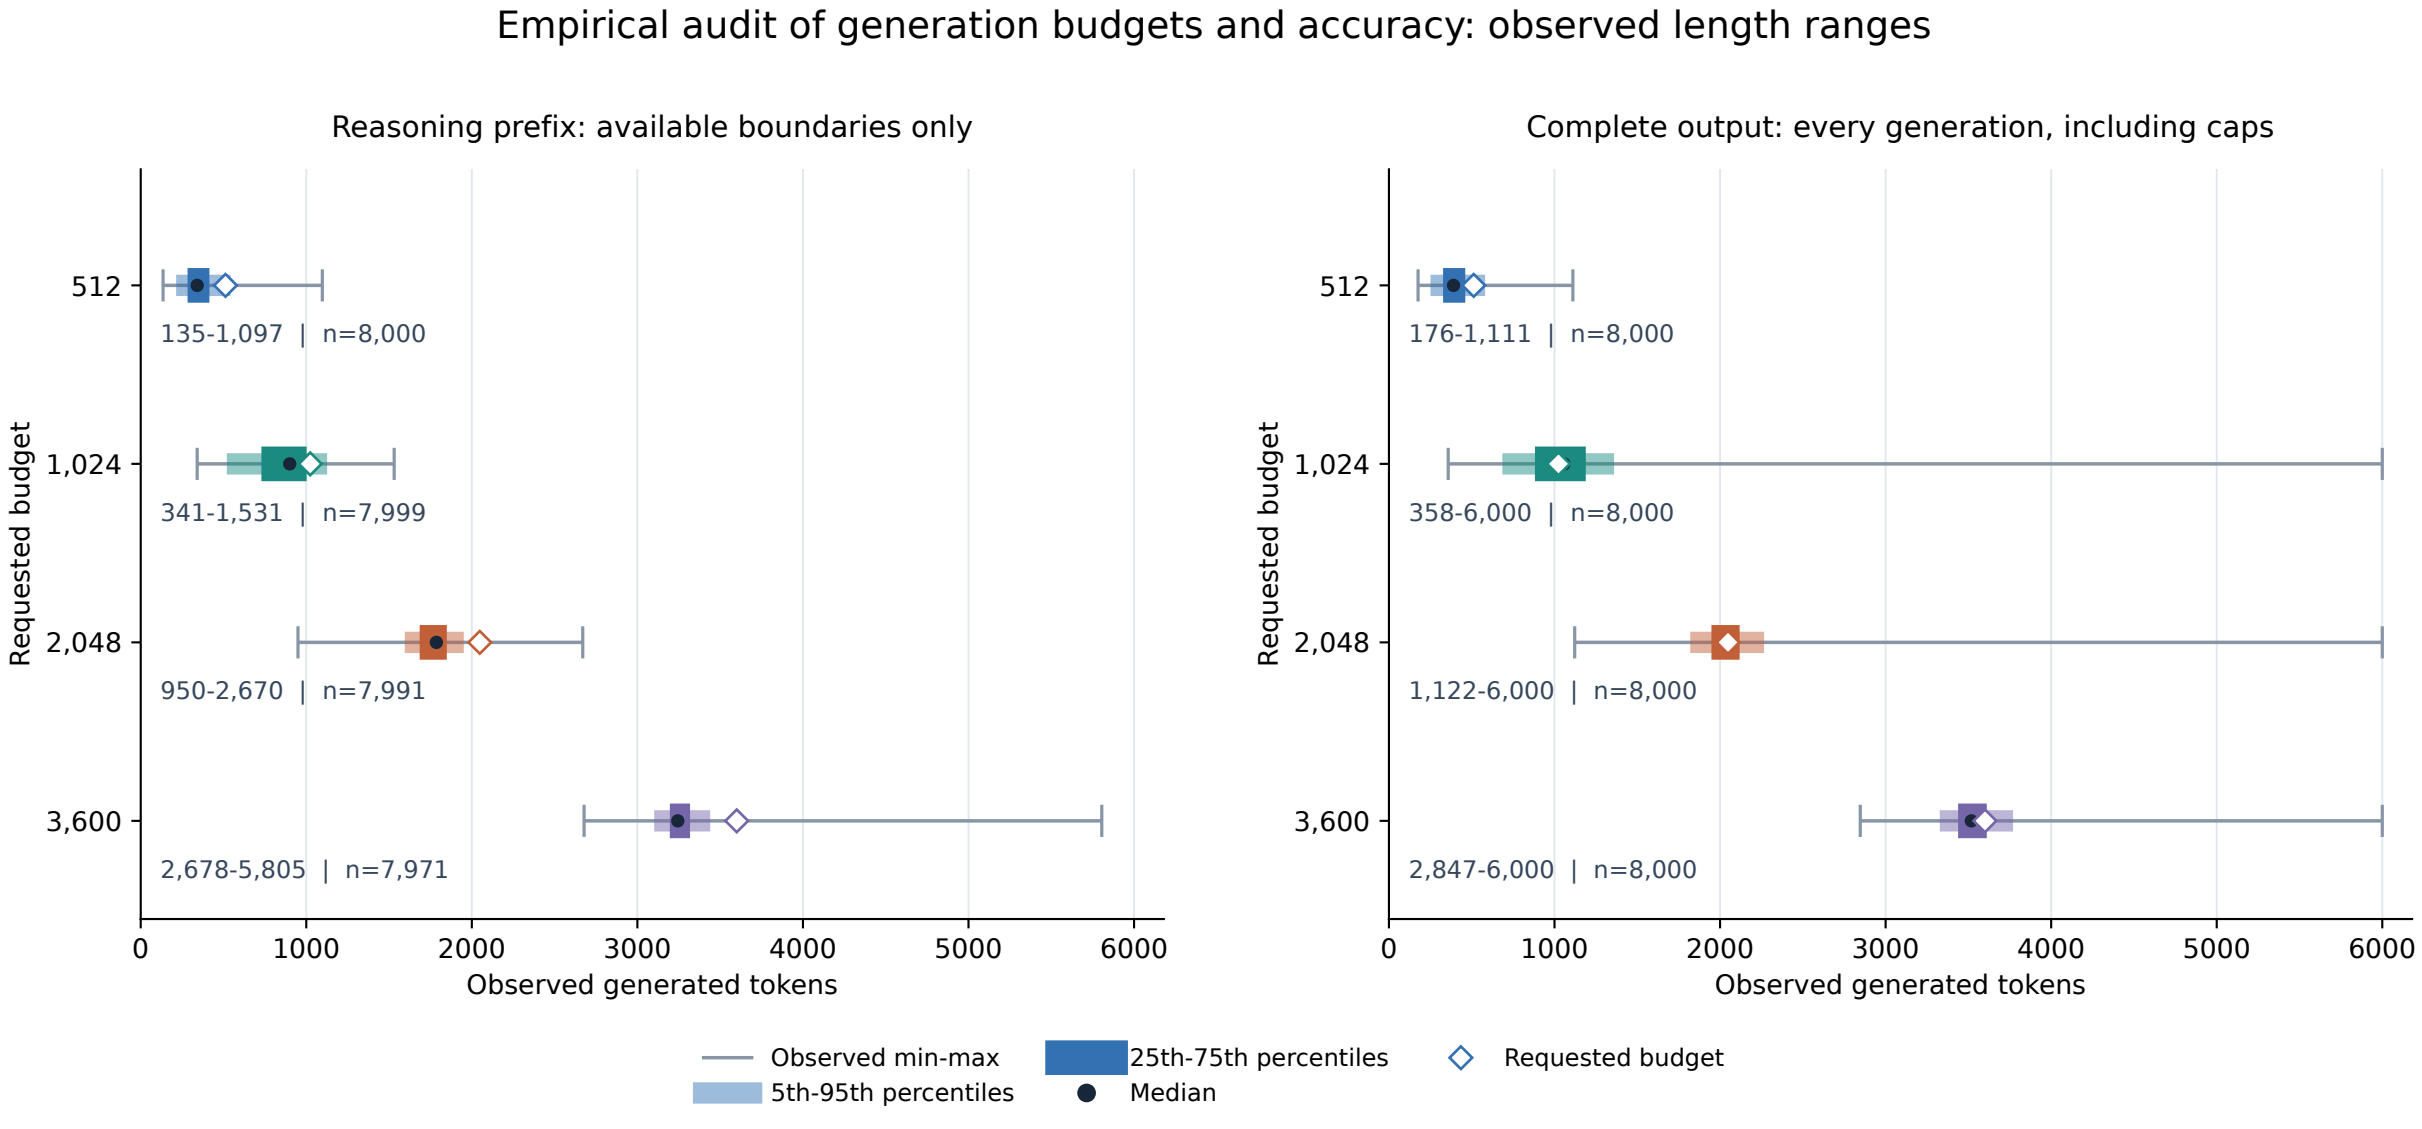
Empirical audit of generation budgets and accuracy: observed length ranges
Reasoning prefix: available boundaries only
Complete output: every generation, including caps
Requested budget
512
1,024
2,048
3,600
135-1,097 | n=8,000
341-1,531 | n=7,999
950-2,670 | n=7,991
2,678-5,805 | n=7,971
176-1,111 | n=8,000
358-6,000 | n=8,000
1,122-6,000 | n=8,000
2,847-6,000 | n=8,000
0
1000
2000
3000
4000
5000
6000
Observed generated tokens
Observed min-max
5th-95th percentiles
25th-75th percentiles
Median
Requested budget
500 questions x 16 generations x 4 budgets. Reasoning lengths unavailable for 39 outputs; total-output cap = 6,000.