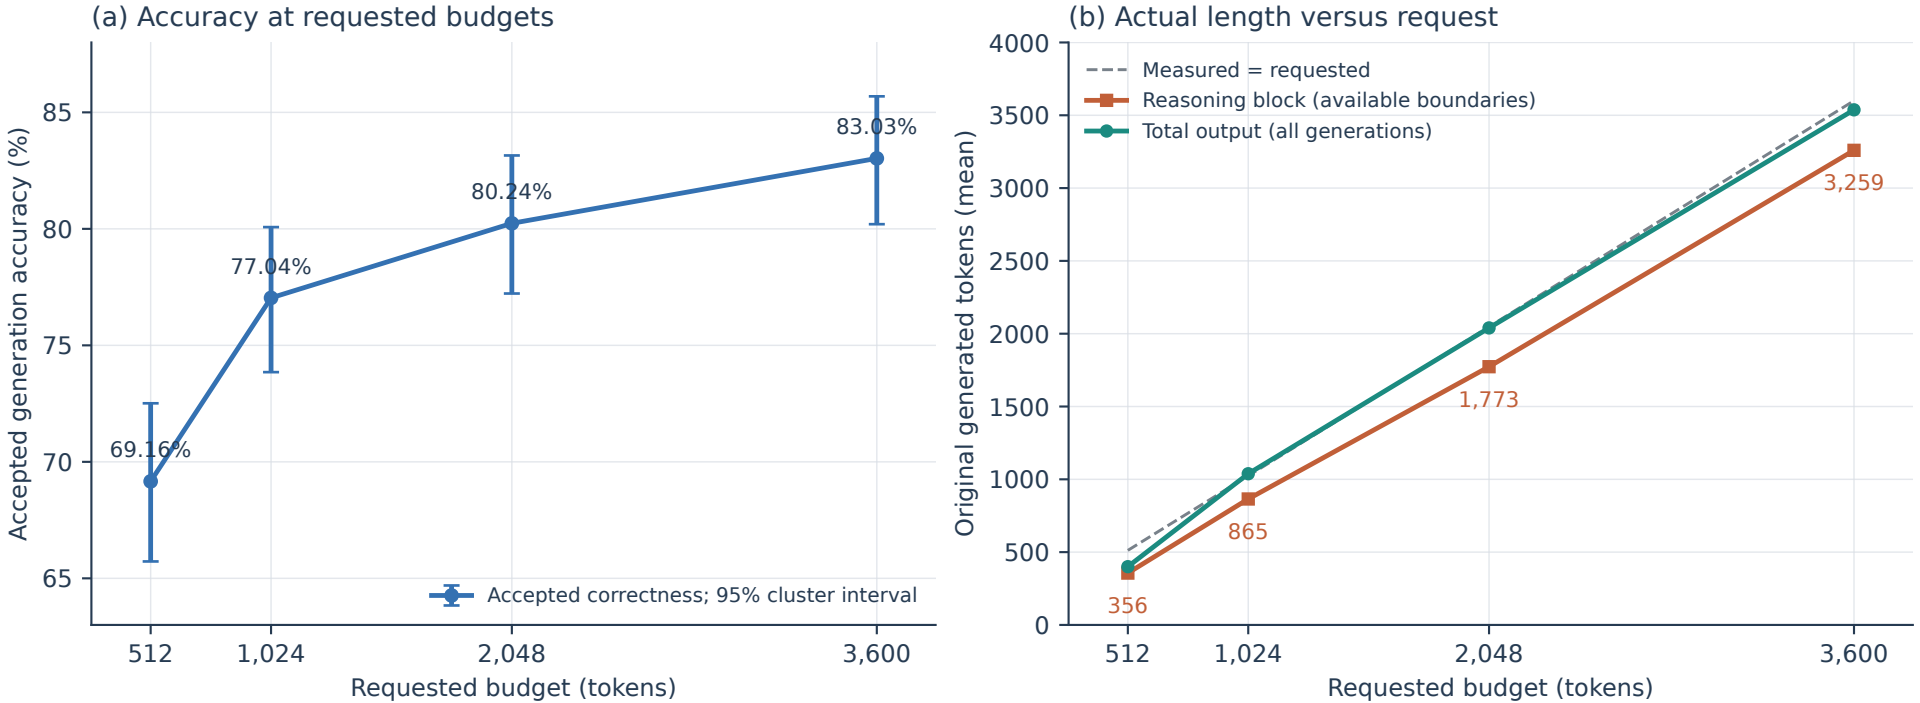
(a) Accuracy at requested budgets
Accepted generation accuracy (%)
85
80
75
70
65
69.16%
77.04%
80.24%
83.03%
Accepted correctness; 95% cluster interval
512
1,024
2,048
3,600
Requested budget (tokens)
(b) Actual length versus request
Original generated tokens (mean)
4000
3500
3000
2500
2000
1500
1000
500
0
Measured = requested
Reasoning block (available boundaries)
Total output (all generations)
356
865
1,773
3,259
512
1,024
2,048
3,600
Requested budget (tokens)